

GENERALIZED ADDITIVE MODELS

Dr. Aric LaBarr

Institute for Advanced Analytics

GENERAL STRUCTURE

Generalized Additive Models (GAMs)

- Provides **general** framework for **adding** of non-linear functions together instead of the typical linear structure.

$$y = \beta_0 + f_1(x_1) + f_2(x_2) + \cdots + f_p(x_p) + \varepsilon$$

- Can be used for regression or classification problems.

Generalized Additive Models (GAMs)

- Provides **general** framework for **adding** of non-linear functions together instead of the typical linear structure.

$$y = \beta_0 + \boxed{f_1(x_1)} + \boxed{f_2(x_2)} + \cdots + \boxed{f_p(x_p)} + \varepsilon$$

- Can be used for regression or classification problems.

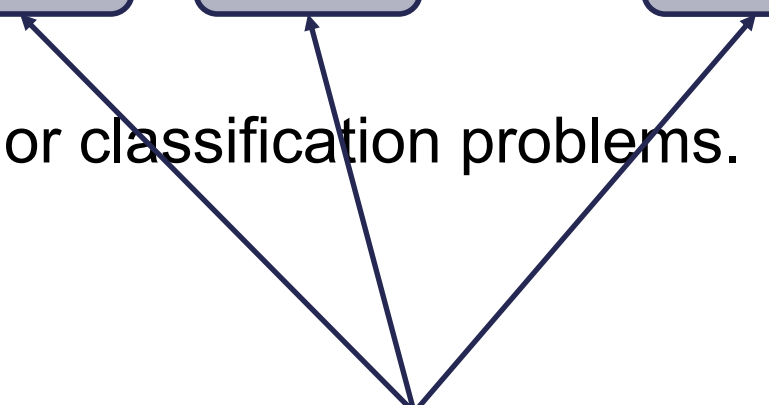
Adding **potentially** complex, individual relationships together.

Generalized Additive Models (GAMs)

- Provides **general** framework for **adding** of non-linear functions together instead of the typical linear structure.

$$y = \beta_0 + \boxed{f_1(x_1)} + \boxed{f_2(x_2)} + \cdots + \boxed{f_p(x_p)} + \varepsilon$$

- Can be used for regression or classification problems.



Many potential complex relationships
to try and model with.



PIECEWISE LINEAR REGRESSION


Changing Slopes

- The slope of the linear relationship between a predictor variable and a response variable can change over different values of the predictor variable.
- Typical straight-line model $\hat{y} = \beta_0 + \beta_1 x_1$ will not be a good fit for this type of data.


Cement Data

- The comprehensive strength of concrete depends on the proportion of water mixed with cement.
- The comprehensive strength decreases at a much faster rate for batches with a greater than 70% water/cement ratio.


Changing Slopes



Changing Slopes



Changing Slopes



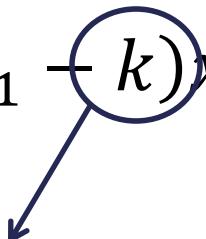
Piecewise Linear Regression

- A model where different straight-line relationships for different intervals in the predictor variable is called the **piecewise linear regression model**.
- The model is the following for two slopes:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k) x_2 + \varepsilon$$

Piecewise Linear Regression

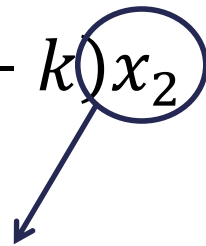
- A model where different straight-line relationships for different intervals in the predictor variable is called the **piecewise linear regression model**.
- The model is the following for two slopes:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k) x_2 + \varepsilon$$


Knot value for x_1 .

Piecewise Linear Regression

- A model where different straight-line relationships for different intervals in the predictor variable is called the **piecewise linear regression model**.
- The model is the following for two slopes:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k)x_2 + \varepsilon$$


$$x_2 = \begin{cases} 1, & x_1 > k \\ 0, & x_1 \leq k \end{cases}$$

Piecewise Linear Regression

- A model where different straight-line relationships for different intervals in the predictor variable is called the **piecewise linear regression model**.
- The model is the following for two slopes:

$$x_2 = 0 \quad \begin{cases} y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k)x_2 + \varepsilon \\ y = \beta_0 + \beta_1 x_1 + \varepsilon \end{cases}$$

Piecewise Linear Regression


- A model where different straight-line relationships for different intervals in the predictor variable is called the **piecewise linear regression model**.
- The model is the following for two slopes:

The diagram illustrates the simplification of a piecewise linear regression equation for two different values of the predictor variable x_2 . A large curved arrow on the left points from the general equation at the top to the two simplified equations below. The top equation is $y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k)x_2 + \varepsilon$. The middle equation, reached by a light blue arrow from $x_2 = 0$, is $y = \beta_0 + \beta_1 x_1 + \varepsilon$. The bottom equation, reached by a dark blue arrow from $x_2 = 1$, is $y = (\beta_0 - k\beta_2) + (\beta_1 + \beta_2)x_1 + \varepsilon$.

$$\begin{array}{l} x_2 = 0 \\ x_2 = 1 \end{array} \quad \begin{array}{l} y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k)x_2 + \varepsilon \\ y = \beta_0 + \beta_1 x_1 + \varepsilon \\ y = (\beta_0 - k\beta_2) + (\beta_1 + \beta_2)x_1 + \varepsilon \end{array}$$

Cement Data

- The comprehensive strength of concrete depends on the proportion of water mixed with cement.
- The comprehensive strength decreases at a much faster rate for batches with a greater than 70% water/cement ratio.



Piecewise Linear Regression

```
cement.lm <- lm(STRENGTH ~ RATIO + X2STAR, data = cement)
```


```
summary(cement.lm)
```

$$(x_1 - k)x_2$$


```
## Call:
## lm(formula = STRENGTH ~ RATIO + X2STAR, data = cement)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.72124 -0.09753 -0.00163  0.24297  0.49393
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   7.79198     0.67696   11.510 7.62e-09 ***
## RATIO         -0.06633     0.01123   -5.904 2.89e-05 ***
## X2STAR        -0.10119     0.02812   -3.598 0.00264 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3286 on 15 degrees of freedom
## Multiple R-squared:  0.9385, Adjusted R-squared:  0.9303
## F-statistic: 114.4 on 2 and 15 DF,  p-value: 8.257e-10
```

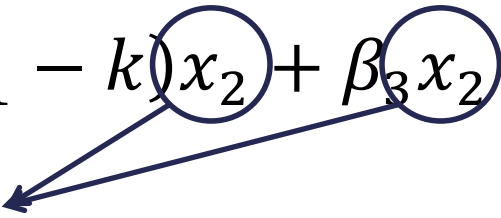
Piecewise Linear Regression

```
ggplot(cement, aes(x = RATIO, y = STRENGTH)) +  
  geom_point() +  
  geom_line(data = cement, aes(x = RATIO,  
                                y = cement.lm$fitted.values)) +  
  ylim(0,6)
```



Extensions – Discontinuous

- The previous approach dealt with piecewise functions that are continuous.
- The following is the discontinuous set-up for two straight lines:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k) x_2 + \beta_3 x_2 + \varepsilon$$


$$x_2 = \begin{cases} 1, & x_1 > k \\ 0, & x_1 \leq k \end{cases}$$

Extensions – Discontinuous

```
cement.lm <- lm(STRENGTH ~ RATIO + X2STAR + X2, data = cement)
```

```
summary(cement.lm)
```

$$(x_1 - k)x_2$$


```
##
## Call:
## lm(formula = STRENGTH ~ RATIO + X2STAR + X2, data = cement)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.53167	-0.15513	0.06171	0.17239	0.49451


```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	7.04975	0.68558	10.283	6.6e-08	***
RATIO	-0.05240	0.01174	-4.463	0.000536	***
X2STAR	-0.07888	0.02686	-2.937	0.010830	*
X2	-0.60388	0.26877	-2.247	0.041302	*

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2916 on 14 degrees of freedom
## Multiple R-squared:  0.9548, Adjusted R-squared:  0.9451
## F-statistic: 98.57 on 3 and 14 DF,  p-value: 1.188e-09
```

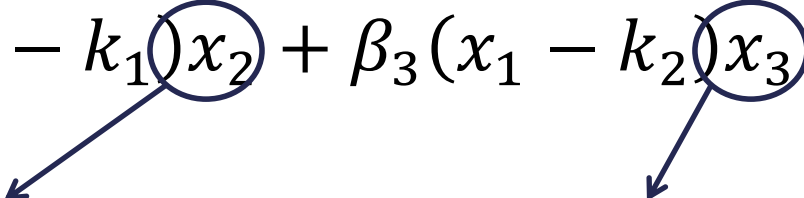
Extensions – Discontinuous

```
qplot(RATIO, STRENGTH, group = X2,  
      geom = c('point', 'smooth'),  
      method = 'lm', data = cement,  
      ylim = c(0,6))
```



Extensions

- The same modeling approach can be applied to any piecewise regression.
- The following is the set-up for three straight lines:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 (x_1 - k_1) x_2 + \beta_3 (x_1 - k_2) x_3 + \varepsilon$$
The diagram shows two arrows originating from the circled terms in the equation above. One arrow points from the circled x_2 to the definition of x_2 below. The other arrow points from the circled x_3 to the definition of x_3 below.

$$x_2 = \begin{cases} 1, & x_1 > k_1 \\ 0, & \text{if not} \end{cases}$$

$$x_3 = \begin{cases} 1, & x_1 > k_2 \\ 0, & \text{if not} \end{cases}$$




MARS (AND EARTH)

Multivariate Adaptive Regression Splines (MARS)

- Multivariate adaptive regression splines (MARS) is a non-parametric technique that is still has a linear form to the model (additive) but has nonlinearities and interaction between variables.
- Essentially, uses **piecewise** regression approach to split into pieces then potentially uses either linear or nonlinear patterns for each piece.


Knots

- MARS first looks for the point in the range of x where two linear functions on either side of the point provides the least squared error (linear regression).




Knots

- MARS first looks for the point in the range of x where two linear functions on either side of the point provides the least squared error (linear regression).




Knots

- MARS first looks for the point in the range of x where two linear functions on either side of the point provides the least squared error (linear regression).

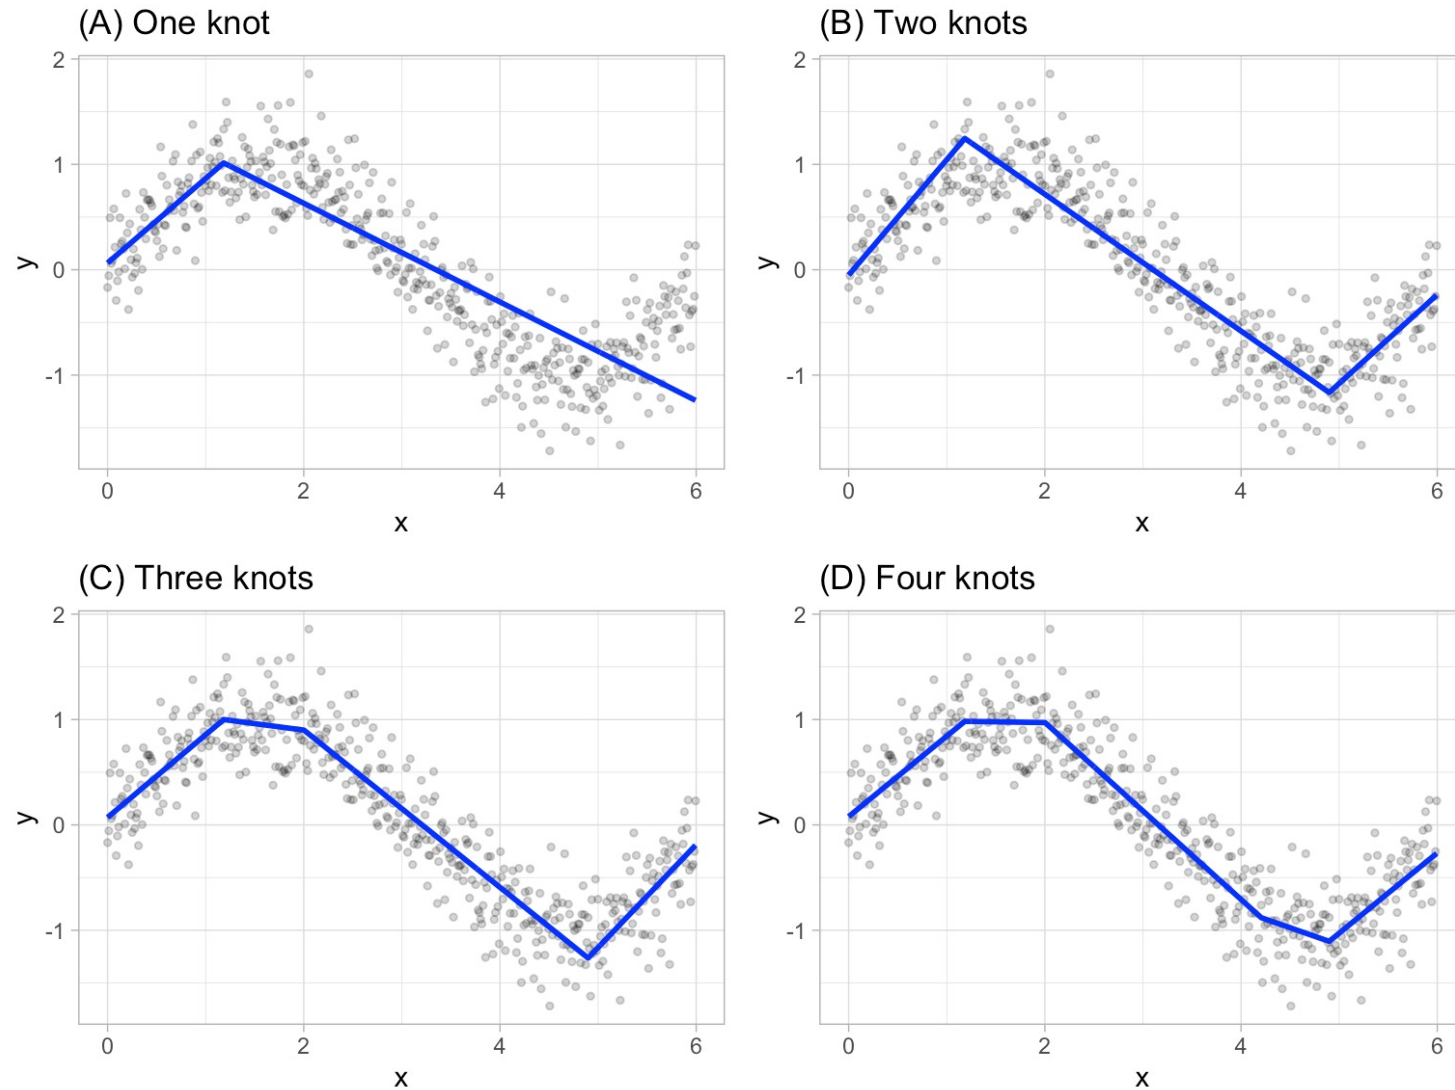


Knots

- MARS first looks for the point in the range of x where two linear functions on either side of the point provides the least squared error (linear regression).



Many Knots



How Many Knots?

- Algorithm continues on each piece of the piecewise function until many knots are found (WILL OVERFIT YOUR DATA).
- Then works backwards (“prunes”) to remove knots that do not contribute significantly to out of sample accuracy.
 - Calculation is performed by the generalized cross-validation (GCV) procedure – computational shortcut for leave-one-out cross-validation.
- **Does this for all variables!**

$$y = \beta_0 + f_1(x_1) + f_2(x_2) + \cdots + f_p(x_p) + \varepsilon$$

EARTH

- **E**nhanced **A**daptive **R**egression **T**hrough **H**inges (EARTH) is the implementation of the MARS algorithm in most software.
- MARS is trademarked by Salford Systems, so instead we use EARTH.

EARTH (and MARS)

```
ames <- make_ordinal_ames()
```

```
ames <- ames %>% mutate(id = row_number())
```

```
set.seed(4321)
```

```
training <- ames %>% sample_frac(0.7)
```


```
testing <- anti_join(ames, training, by = 'id')
```

```
training <- training %>%
```

```
  select(Sale_Price, Bedroom_AbvGr, Year_Built, Mo_Sold, Lot_Area, Street,  
         Central_Air, First_Flr_SF, Second_Flr_SF, Full_Bath, Half_Bath, Fireplaces,  
         Garage_Area, Gr_Liv_Area, TotRms_AbvGrd)
```

EARTH (and MARS)

```
ggplot(training, aes(x = Garage_Area,  
                     y = Sale_Price)) +  
  geom_point()
```




EARTH (and MARS)

```
mars1 <- earth(Sale_Price ~ Garage_Area, data = training)
summary(mars1)
```

```
## Call: earth(formula=Sale_Price~Garage_Area, data=training)
##
##               coefficients
## (Intercept)      124159.039
## h(286-Garage_Area)    -60.257
## h(Garage_Area-286)    297.277
## h(Garage_Area-521)   -483.642
## h(Garage_Area-576)    733.859
## h(Garage_Area-758)   -356.460
## h(Garage_Area-1043)  -490.873
##
## Selected 7 of 7 terms, and 1 of 1 predictors
## Termination condition: RSq changed by less than 0.001 at 7 terms
## Importance: Garage_Area
## Number of terms at each degree of interaction: 1 6 (additive model)
## GCV 3427475346    RSS 6.94092e+12    GRSq 0.4492014    RSq 0.4556309
```

EARTH (and MARS)

```
ggplot(training, aes(x = Garage_Area,  
                     y = Sale_Price)) +  
  geom_point() +  
  geom_line(data = training,  
            aes(x = Garage_Area,  
                y = mars1$fitted.values),  
            color = "blue")
```



EARTH (and MARS)

```

mars1 <- earth(Sale_Price ~ ., data = training)
summary(mars1)

## Call: earth(formula=Sale_Price~., data=training)
##
##               coefficients
## (Intercept)      319493.46
## Central_AirY      20289.49
## h(4-Bedroom_AbvGr)  9214.66
## h(Bedroom_AbvGr-4) -23009.05
## ...
## h(Half_Bath-1)      -45378.31
## h(2-Fireplaces)     -14408.56
## h(Fireplaces-2)     -26072.58
## h(Garage_Area-539)   101.97
## h(Garage_Area-1043) -294.30
## h(Gr_Liv_Area-2049)  65.21
## h(Gr_Liv_Area-3194) -159.79
##
## Selected 21 of 24 terms, and 10 of 14 predictors
## Termination condition: Reached nk 29
## Importance: First_Flr_SF, Second_Flr_SF, Year_Built, Garage_Area, ...
## Number of terms at each degree of interaction: 1 20 (additive model)
## GCV 1033819964    RSS 2.036439e+12    GRSq 0.8338641    RSq 0.8402842

```

Variable Importance

```

mars1 <- earth(Sale_Price ~ ., data = training)
summary(mars1)

## Call: earth(formula=Sale_Price~., data=training)
##
##               coefficients
## (Intercept)      319493.46
## Central_AirY      20289.49
## h(4-Bedroom_AbvGr)  9214.66
## h(Bedroom_AbvGr-4) -23009.05
## ...
## h(Half_Bath-1)      -45378.31
## h(2-Fireplaces)     -14408.56
## h(Fireplaces-2)     -26072.58
## h(Garage_Area-539)   101.97
## h(Garage_Area-1043) -294.30
## h(Gr_Liv_Area-2049)  65.21
## h(Gr_Liv_Area-3194) -159.79
##
## Selected 21 of 24 terms, and 10 of 14 predictors
## Termination condition: Reached nk 29
## Importance: First Flr SF, Second Flr SF, Year Built, Garage Area, ...
## Number of terms at each degree of interaction: 1 20 (additive model)
## GCV 1033819964    RSS 2.036439e+12    GRSq 0.8338641    RSq 0.8402842

```


Variable Importance

- There is one “subset” for each model size (1 term, 2 terms, etc.) – the best model of that size.
- Ranks variables by how many of the subsets that variable appears in.
 - More subsets of models it appears in (in the best 1 variable model, in the best 2 variable model, etc.), then the better the variable.
- RSS (residual sum of squares or sum of squares error) is scaled version of decrease in residual sum of squares relative to the previous subset.
- GCV is approximation of RSS on leave-one-out-cross validation.


Variable Importance

```
evimp(mars1)
```

##	nsubsets	gcv	rss
## First_Flr_SF	20	100.0	100.0
## Second_Flr_SF	19	71.7	71.9
## Year_Built	18	50.9	51.3
## Garage_Area	17	34.3	35.0
## Fireplaces	16	31.0	31.7
## Gr_Liv_Area	15	27.6	28.4
## Central_AirY	12	20.0	20.9
## Bedroom_AbvGr	11	18.1	19.0
## Lot_Area	10	16.2	17.2
## Half_Bath	4	7.4	8.2

Interpretability

- Can view the “relationship” between predictors and target variable.





SMOOTHING

Smoothing

- GAMs can be made up of any non-parametric function of the predictor variables.
- Another popular technique is to use **smoothing functions** so the piecewise linear regressions are not so jagged.
- Many different types of smoothing functions:
 - LOESS (localized regression)
 - Smoothing splines
 - Regression splines


Smoothing

- GAMs can be made up of any non-parametric function of the predictor variables.
- Another popular technique is to use **smoothing functions** so the piecewise linear regressions are not so jagged.
- Many different types of smoothing functions:
 - LOESS (localized regression)
 - Smoothing splines
 - Regression splines


LOESS

- Locally estimated scatterplot smoothing (LOESS) is a popular smoothing technique.
- Same technique used in STL decomposition in time series


U.S. AIRLINE PASSENGERS
(SEASONALLY ADJUSTED)




LOESS




LOESS




U.S. AIRLINE PASSENGERS (SEASONALLY ADJUSTED)




LOESS




U.S. AIRLINE PASSENGERS (SEASONALLY ADJUSTED)




LOESS




U.S. AIRLINE PASSENGERS (SEASONALLY ADJUSTED)



LOESS



U.S. AIRLINE PASSENGERS (SEASONALLY ADJUSTED)




Smoothing

- GAMs can be made up of any non-parametric function of the predictor variables.
- Another popular technique is to use **smoothing functions** so the piecewise linear regressions are not so jagged.
- Many different types of smoothing functions:
 - LOESS (localized regression)
 - Smoothing splines
 - Regression splines

Smoothing Splines

- Smoothing splines take a different approach as compared to LOESS.
- Smoothing splines have a knot **at every single observation** for piecewise regression – OVERFITTING!
- Use penalty parameter to counterbalance the “wiggle” of the spline.



Smoothing Splines

- Smoothing splines try to find the function $s(x_i)$ that optimally fits x to the target variable y through this equation:

$$\min \sum_{i=1}^n (y_i - s(x_i))^2 + \lambda \int s''(t_i)^2 dt$$

Smoothing Splines

- Smoothing splines try to find the function $s(x_i)$ that optimally fits x to the target variable y through this equation:

$$\min \sum_{i=1}^n (y_i - s(x_i))^2 + \lambda \int s''(t_i)^2 dt$$

Sum of squared error!

Smoothing Splines

- Smoothing splines try to find the function $s(x_i)$ that optimally fits x to the target variable y through this equation:

$$\min \sum_{i=1}^n (y_i - s(x_i))^2 + \lambda \int s''(t_i)^2 dt$$

Penalty (λ) applied to
integral of second
derivative of smoothing
function

Smoothing Splines

- Smoothing splines try to find the function $s(x_i)$ that optimally fits x to the target variable y through this equation:

$$\min \sum_{i=1}^n (y_i - s(x_i))^2 + \lambda \int s''(t_i)^2 dt$$

Penalty (λ) applied to integral of “slope of slopes” which is large when lots of “wiggle”

Smoothing Splines

- Smoothing splines try to find the function $s(x_i)$ that optimally fits x to the target variable y through this equation:

$$\min \sum_{i=1}^n (y_i - s(x_i))^2 + \lambda \int s''(t_i)^2 dt$$

Penalty (λ) estimated with another approximation of leave one out cross validation

Smoothing

- GAMs can be made up of any non-parametric function of the predictor variables.
- Another popular technique is to use **smoothing functions** so the piecewise linear regressions are not so jagged.
- Many different types of smoothing functions:
 - LOESS (localized regression)
 - Smoothing splines
 - Regression splines – computationally nicer version of smoothing splines


GAM's with Splines

```
gam1 <- mgcv::gam(Sale_Price ~ s(Garage_Area), data = training)
summary(gam1)
```

```
## Family: gaussian
## Link function: identity
##
## Formula:
## Sale_Price ~ s(Garage_Area)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   180897      1290    140.2  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(Garage_Area) 8.134  8.769 192.5  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.451   Deviance explained = 45.3%
## GCV = 3.4301e+09   Scale est. = 3.4148e+09   n = 2051
```

GAM's with Splines

```
plot(gam1)
```



GAM's with Splines

```
gam2 <- mgcv::gam(Sale_Price ~ s(Bedroom_AbvGr, k = 5) +  
  s(Year_Built) +  
  s(Mo_Sold) +  
  s(Lot_Area) +  
  s(First_Flr_SF) +  
  s(Second_Flr_SF) +  
  s(Garage_Area) +  
  s(Gr_Liv_Area) +  
  s(TotRms_AbvGrd) +  
  Street +  
  Central_Air +  
  factor(Fireplaces) +  
  factor(Full_Bath) +  
  factor(Half_Bath)  
  , method = 'REML', data = training)  
summary(gam2)
```


GAM's with Splines

```
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    136139     19681   6.917 6.19e-12 ***
## StreetPave      27689      12710   2.178  0.0295  *
## Central_AirY    18012       3168   5.685 1.50e-08 ***
## factor(Fireplaces)1  14070       1666   8.443 < 2e-16 ***
## factor(Fireplaces)2  27137       3146   8.626 < 2e-16 ***
## factor(Fireplaces)3  15704      10552   1.488  0.1368
## factor(Fireplaces)4 -79595      31469  -2.529  0.0115  *
## factor(Full_Bath)1   -5341      14528  -0.368  0.7132
## factor(Full_Bath)2  -11074      14827  -0.747  0.4552
## factor(Full_Bath)3    1226      15787   0.078  0.9381
## factor(Full_Bath)4  -16271      24326  -0.669  0.5037
## factor(Half_Bath)1    2102       2206   0.953  0.3408
## factor(Half_Bath)2  -38507       9111  -4.226 2.48e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
...

```

GAM's with Splines

```
...
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(Bedroom_AbvGr) 2.653  3.165  18.789 <2e-16 ***
## s(Year_Built)    6.445  7.543 101.758 <2e-16 ***
## s(Mo_Sold)       1.516  1.868   0.993  0.4507
## s(Lot_Area)      7.186  8.193  11.726 <2e-16 ***
## s(First_Flr_SF)  8.063  8.765  15.548 <2e-16 ***
## s(Second_Flr_SF) 8.212  8.818   7.806 <2e-16 ***
## s(Garage_Area)   7.426  8.328  21.654 <2e-16 ***
## s(Gr_Liv_Area)   8.545  8.882  14.834 <2e-16 ***
## s(TotRms_AbvGrd) 3.805  4.738   1.921  0.0783 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) = 0.851  Deviance explained = 85.6%
## -REML = 23957  Scale est. = 9.2392e+08  n = 2051
```

GAM's with Selected Splines

```
sel.gam2 <- mgcv::gam(Sale_Price ~ s(Bedroom_AbvGr, k = 5) +
  s(Year_Built) +
  s(Mo_Sold) +
  s(Lot_Area) +
  s(First_Flr_SF) +
  s(Second_Flr_SF) +
  s(Garage_Area) +
  s(Gr_Liv_Area) +
  s(TotRms_AbvGrd) +
  Street +
  Central_Air +
  factor(Fireplaces) +
  factor(Full_Bath) +
  factor(Half_Bath)
, method = 'REML',
  select = TRUE, data = training)

summary(sel.gam2)
```

```
## Approximate significance of smooth terms:
##               edf Ref.df      F  p-value
## s(Bedroom_AbvGr) 2.333381     4 14.833 < 2e-16 ***
## s(Year_Built)    6.945994     9 83.962 < 2e-16 ***
## s(Mo_Sold)       0.007522     9  0.001 0.33142
## s(Lot_Area)      7.561273     9 11.790 < 2e-16 ***
## s(First_Flr_SF)  8.602695     9 32.390 < 2e-16 ***
## s(Second_Flr_SF) 0.940878     9  1.716 3.89e-06 ***
## s(Garage_Area)   6.688676     9 19.626 < 2e-16 ***
## s(Gr_Liv_Area)   4.839941     9  7.634 < 2e-16 ***
## s(TotRms_AbvGrd) 3.740058     9  1.280 0.00845 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.845   Deviance explained = 84.9%
## -REML = 24081   Scale est. = 9.6657e+08   n = 2051
```



SUMMARY

Generalized Additive Models (GAMs)

- Provides **general** framework for **adding** of non-linear functions together instead of the typical linear structure.

$$y = \beta_0 + f_1(x_1) + f_2(x_2) + \cdots + f_p(x_p) + \varepsilon$$

- Can be used for regression or classification problems.

Generalized Additive Models (GAMs)

Advantages

- Allows a nonlinear relationship without trying out many transformations manually
- Improved predictions
- Still has some “interpretation”
- Computationally fast

Disadvantages

- Can incorporate interactions but can take time
- Not good for large numbers of variables – **prescreening needed!**
- Multicollinearity still a problem

